

Cloudflare TinyURL - Design Document

Overview

Cloudflare TinyURL is a URL shortening service designed to efficiently generate short URLs, track their usage, and ensure data persistence with high scalability and fault tolerance.

Requirements

Functional Requirements

1. **Create Short URL:** Ensure uniqueness and prevent duplicates.
2. **Redirect Short URL:** Retrieve the original URL and redirect the request.
3. **Click Tracking:**
 - Track the number of clicks in different time frames: last 24 hours, last week, all time, and last minute.
 - Provide a fallback mechanism to query PostgreSQL when Redis is unavailable.
4. **Persistence:**
 - Ensure all URLs and click events are stored permanently in PostgreSQL.
 - Maintain Redis for fast lookup with persistence enabled.
5. **Delete Short URL:** Remove URLs upon request.
6. **Metrics and Logging:** Provide observability into system operations (optional but recommended).

Non-Functional Requirements

- **Scalability:** The system should be horizontally scalable.
- **Caching:** Use Redis to cache URLs for low-latency responses.
- **Event-Driven Architecture:** Process click events asynchronously.
- **Fault Tolerance:** Fallback to PostgreSQL when Redis is unavailable.
- **Testing:** End-to-end system tests to ensure reliability.

API Endpoints

Method	Endpoint	Description
POST	<code>/api/v1/create</code>	Create a short URL
GET	<code>/api/v1/{shortURL}</code>	Redirect to the original long URL

DELETE	<code>/api/v1/{shortURL}</code>	Delete a short URL
GET	<code>/api/v1/clicks/{shortURL}</code>	Retrieve click statistics from Redis
GET	<code>/api/v1/clicks_fallback/{shortURL}</code>	Retrieve click statistics from PostgreSQL (fallback)

System Architecture

Data Entities

PostgreSQL

1. **urls Table**: Stores URL mappings
 - `short_url_id` (Primary Key)
 - `long_url`
 - `created_time`
 - `expire_time`
2. **clicks_urls Table**: Stores all click events for fault tolerance
 - `id` (Primary Key)
 - `short_url`
 - `accessed_at`

Redis

1. **Global Counter**: Used to generate unique short URLs.
2. **Per URL Click Counters**:
 - `last_24hrs`, `last_week`, `all_time`, `last_1min` (with TTL)
3. **Redis Queue**:
 - Stores click events that expire to decrement counts asynchronously.

Key Design Considerations

1. **Caching for URL Redirection**:
 - Instead of querying PostgreSQL on every redirect, we cache URL mappings in Redis.
2. **Click Tracking via Redis Counters**:

- Counting clicks in real-time using Redis counters.
 - Periodically persisting data in PostgreSQL for reliability.
 - Using Redis TTL key expiry events to queue updates for global counters.
3. **Fault Tolerance:**
- If Redis goes down, click statistics can still be retrieved from PostgreSQL.
4. **Event-Based Updates:**
- Redis queue is used to track click expiration events.
 - Events are processed asynchronously to decrement click counts.
5. **Distributed System Considerations:**
- **Short URL Generation:** Base62 encoding + global counter.
 - **Click Events:** Use Snowflake ID for uniqueness in distributed environments.
 - **Horizontal Scalability:** The system can run multiple instances.
-

Testing Strategy


End-to-End Tests

1. **URL Uniqueness Test:** Create 10 URLs and ensure uniqueness.
 2. **URL Deletion Test:** Create and delete 5 URLs, then verify they are removed.
 3. **URL Redirection Test:** Create 10 URLs and validate their redirects.
 4. **Click Counting Test:**
 - Create a URL and redirect it multiple times.
 - Validate click counts at different time intervals.
 - Ensure last-minute counts expire dynamically.
-

Future Enhancements

1. **Improved Logging and Monitoring:** Add distributed tracing and structured logs.
 2. **Scalability Enhancements:** Perform load testing and optimize query performance.
 3. **Advanced Event Processing:** Implement streaming solutions (e.g., Kafka) for click tracking.
 4. **Security Enhancements:** Implement rate-limiting and authentication.
-

System Design Diagram



This design ensures efficient URL handling, caching, event-driven updates, and fault tolerance, making Cloudflare TinyURL a scalable and reliable system.